Training and Validation Performance Across Training Epochs

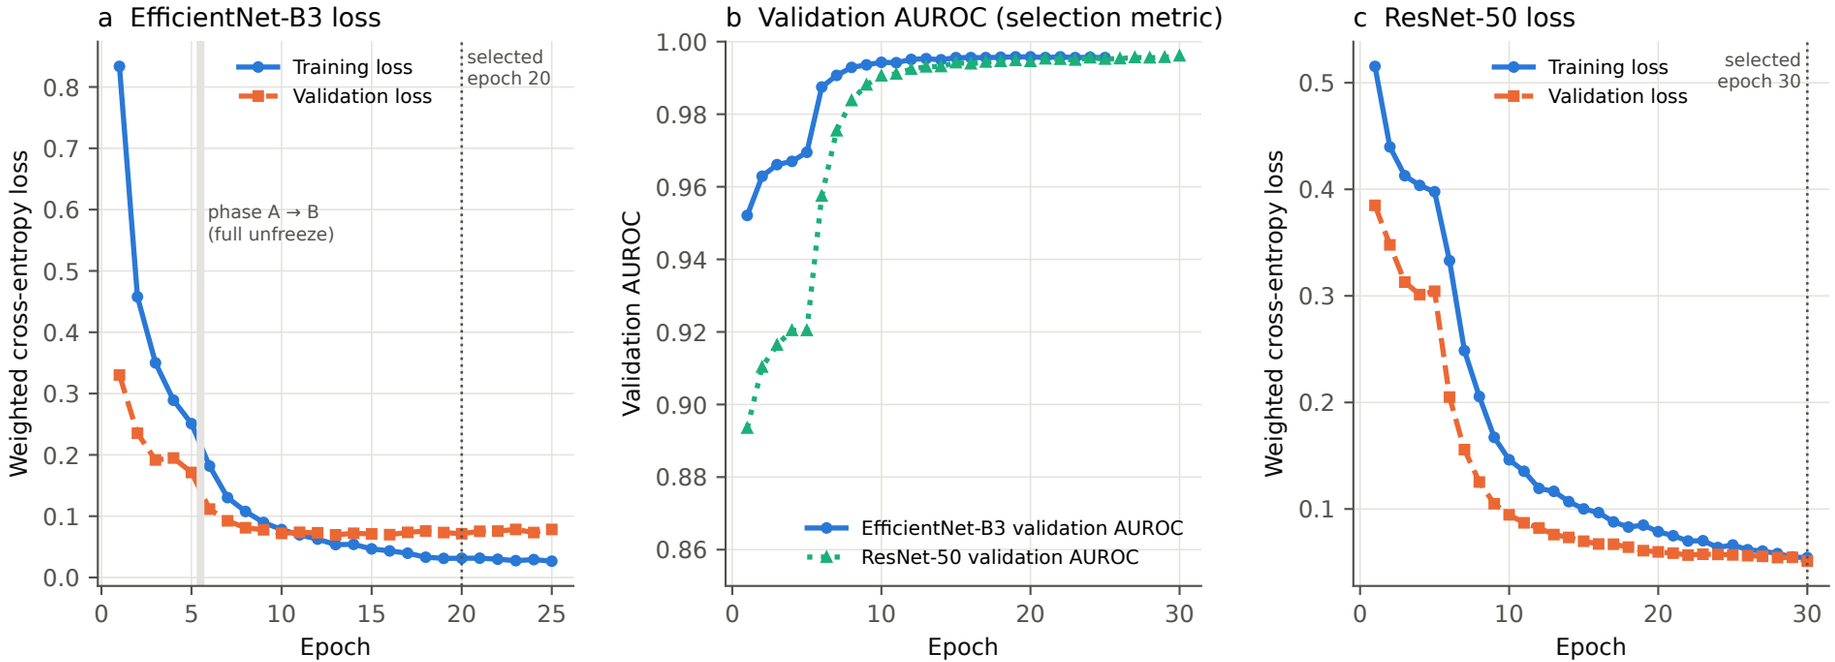

B3 per-epoch values were parsed from the Kaggle console log (artifacts/training_log_reconstructed.csv; the run's training_log.csv was lost); accuracy was not logged for that run. ResNet-50 values from its training_log.csv.
EfficientNet-B0 and the label-smoothing run are not shown (their per-epoch values survive only in console logs, not parsed here). Training loss is computed on augmented images with class weights; validation = patient-disjoint split (n=5,560 after exclusion). Test data never used.