Training and Validation Performance Across Training Epochs

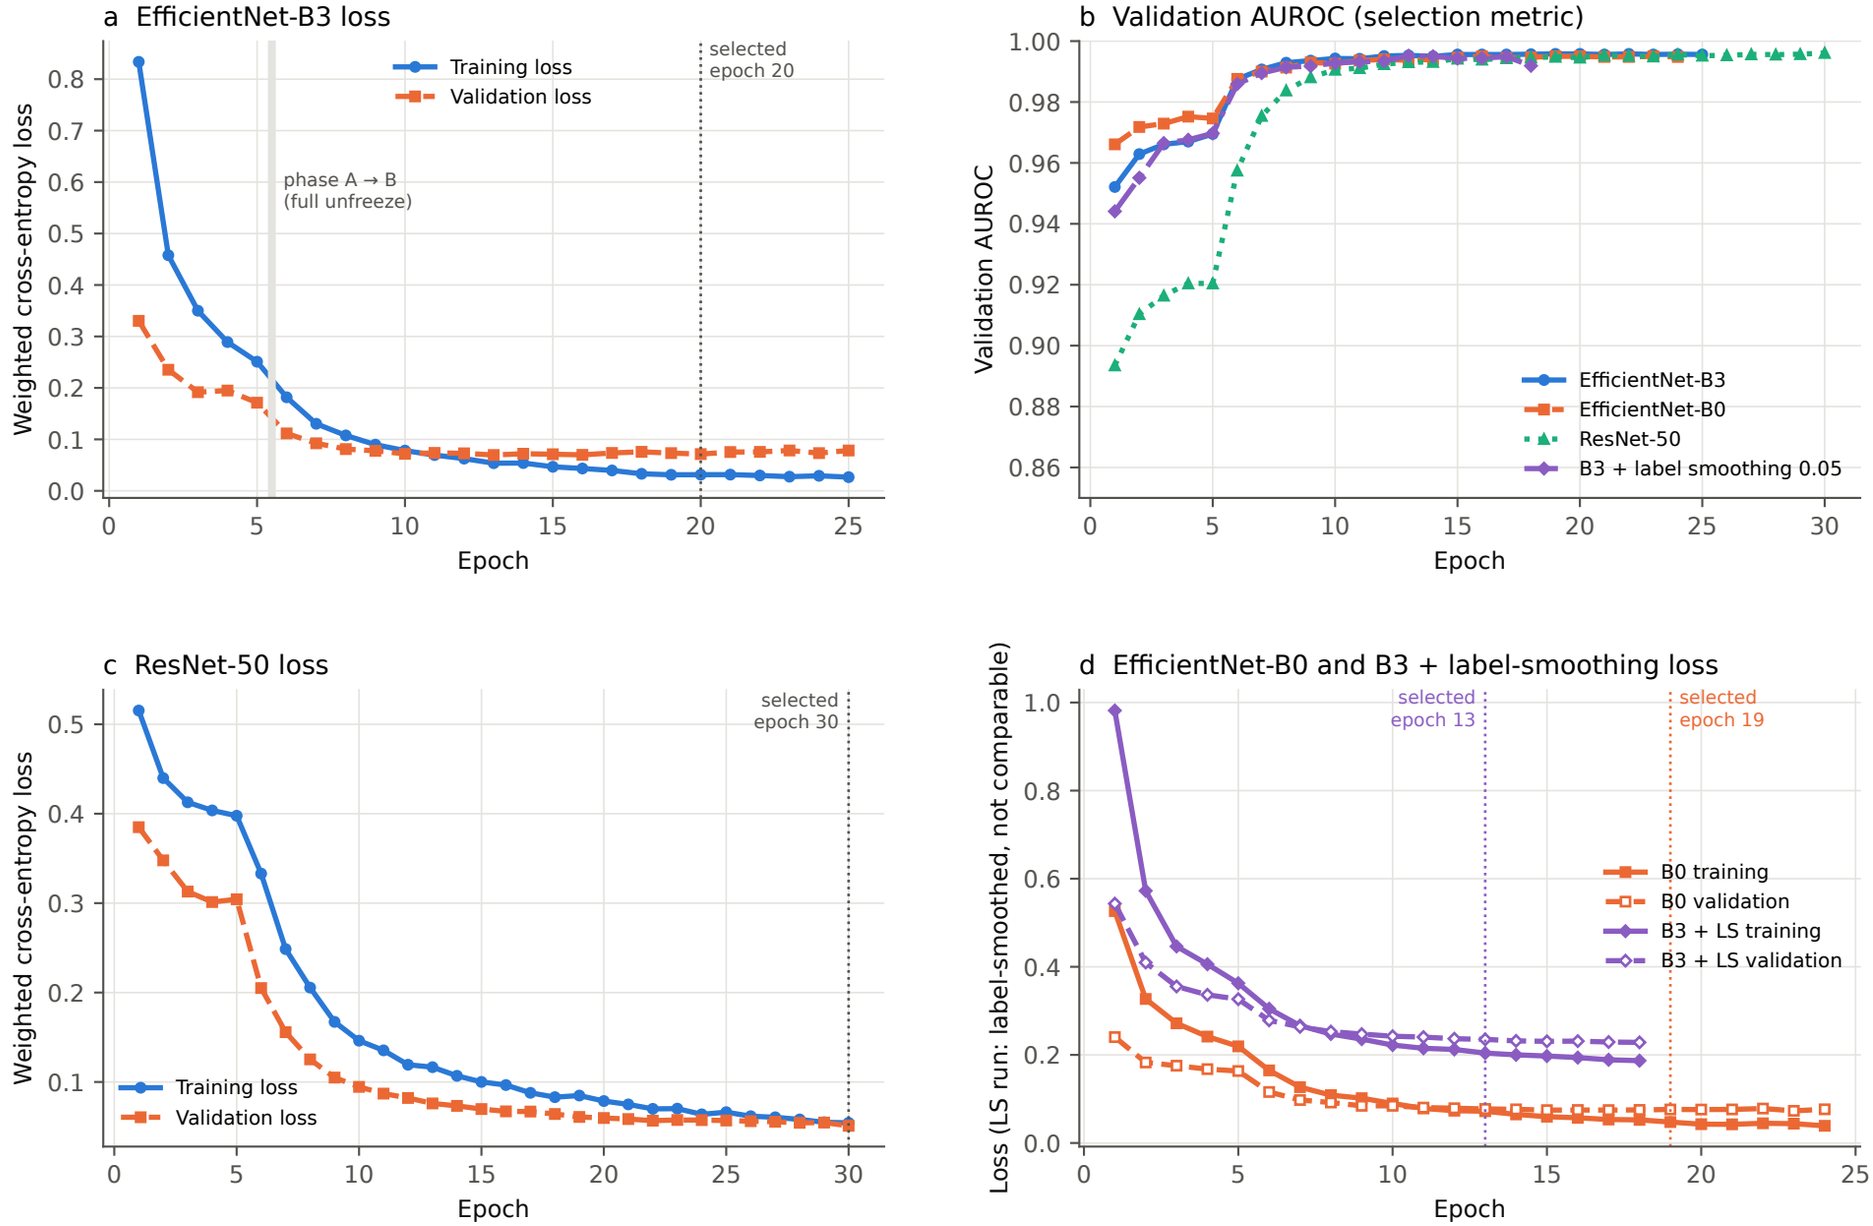

B3, B0 and label-smoothing (LS) per-epoch values were parsed from the Kaggle console logs (B3: artifacts/training_log_reconstructed.csv; B0: kaggle/20260920T095427Z/run.log; LS: kaggle/20260922T160023Z/run.log); accuracy was not logged for those runs. ResNet-50 values from its training_log.csv. Both early-stopped runs end before 30 epochs. The LS loss includes the smoothing target, so its floor is higher and it is not on the same scale as the other losses. Training loss is computed on augmented images with class weights; validation = patient-disjoint split (n=5,560 after exclusion). Test data never used.