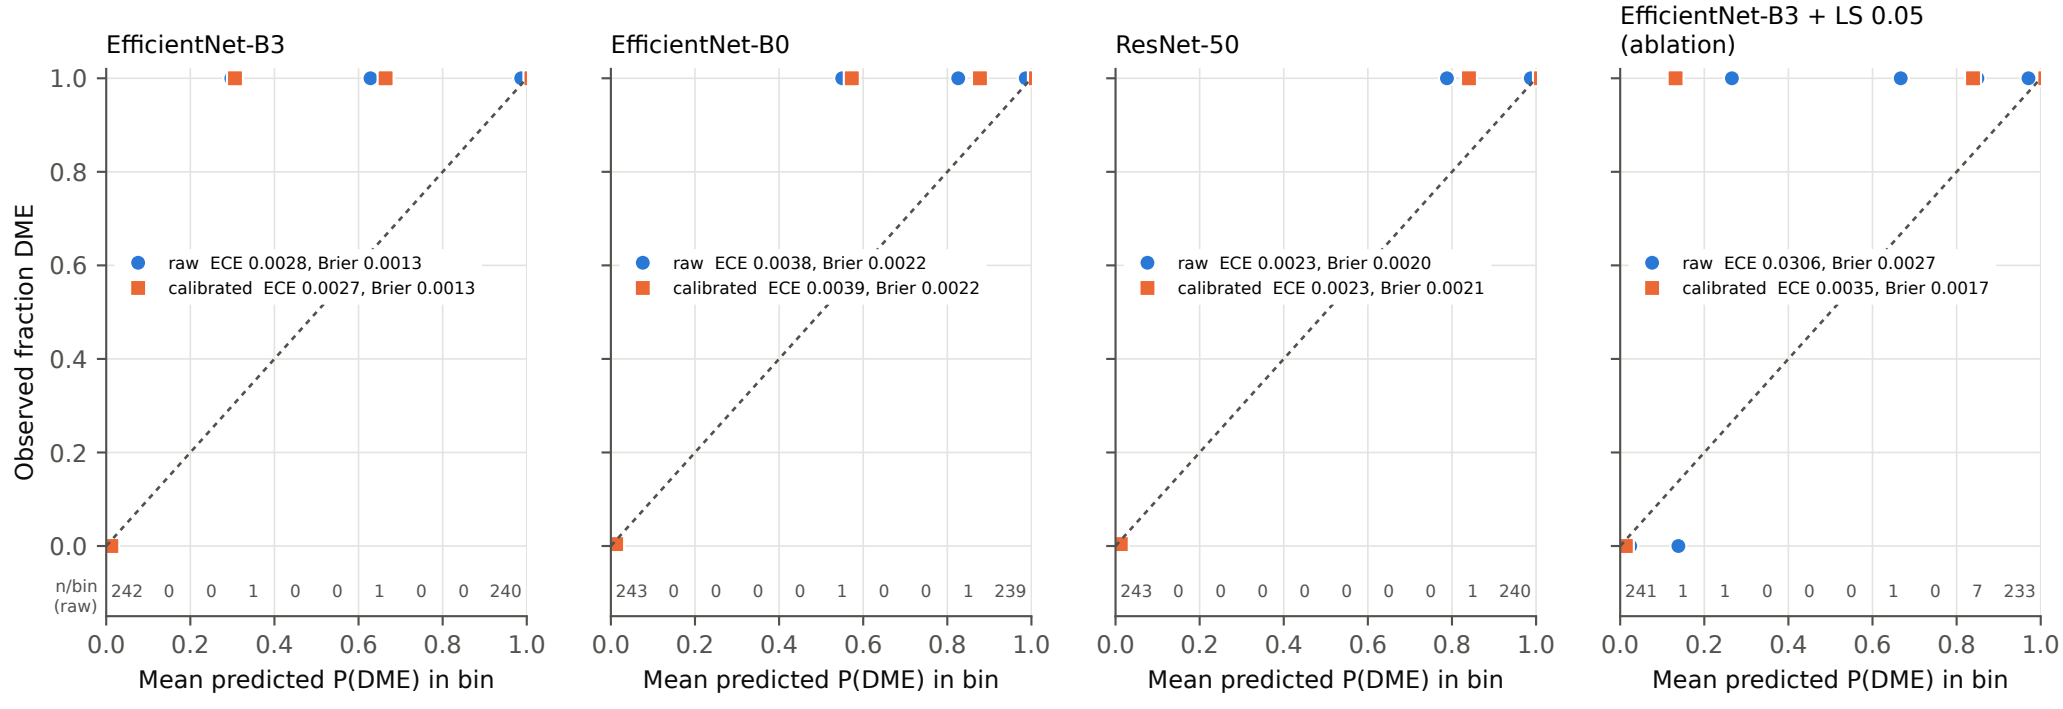

10 equal-width bins (same as evaluate.py); empty bins omitted; numbers under each axis = images per bin (raw). Almost all predictions fall in the outermost bins, so intermediate-bin estimates rest on 0–2 images.
Temperatures were fitted by the original evaluate.py on a re-generated 'validation' split that does not match the training-time split and overlaps the training data (see leakage_audit.md §5); the test set was not used for fitting.