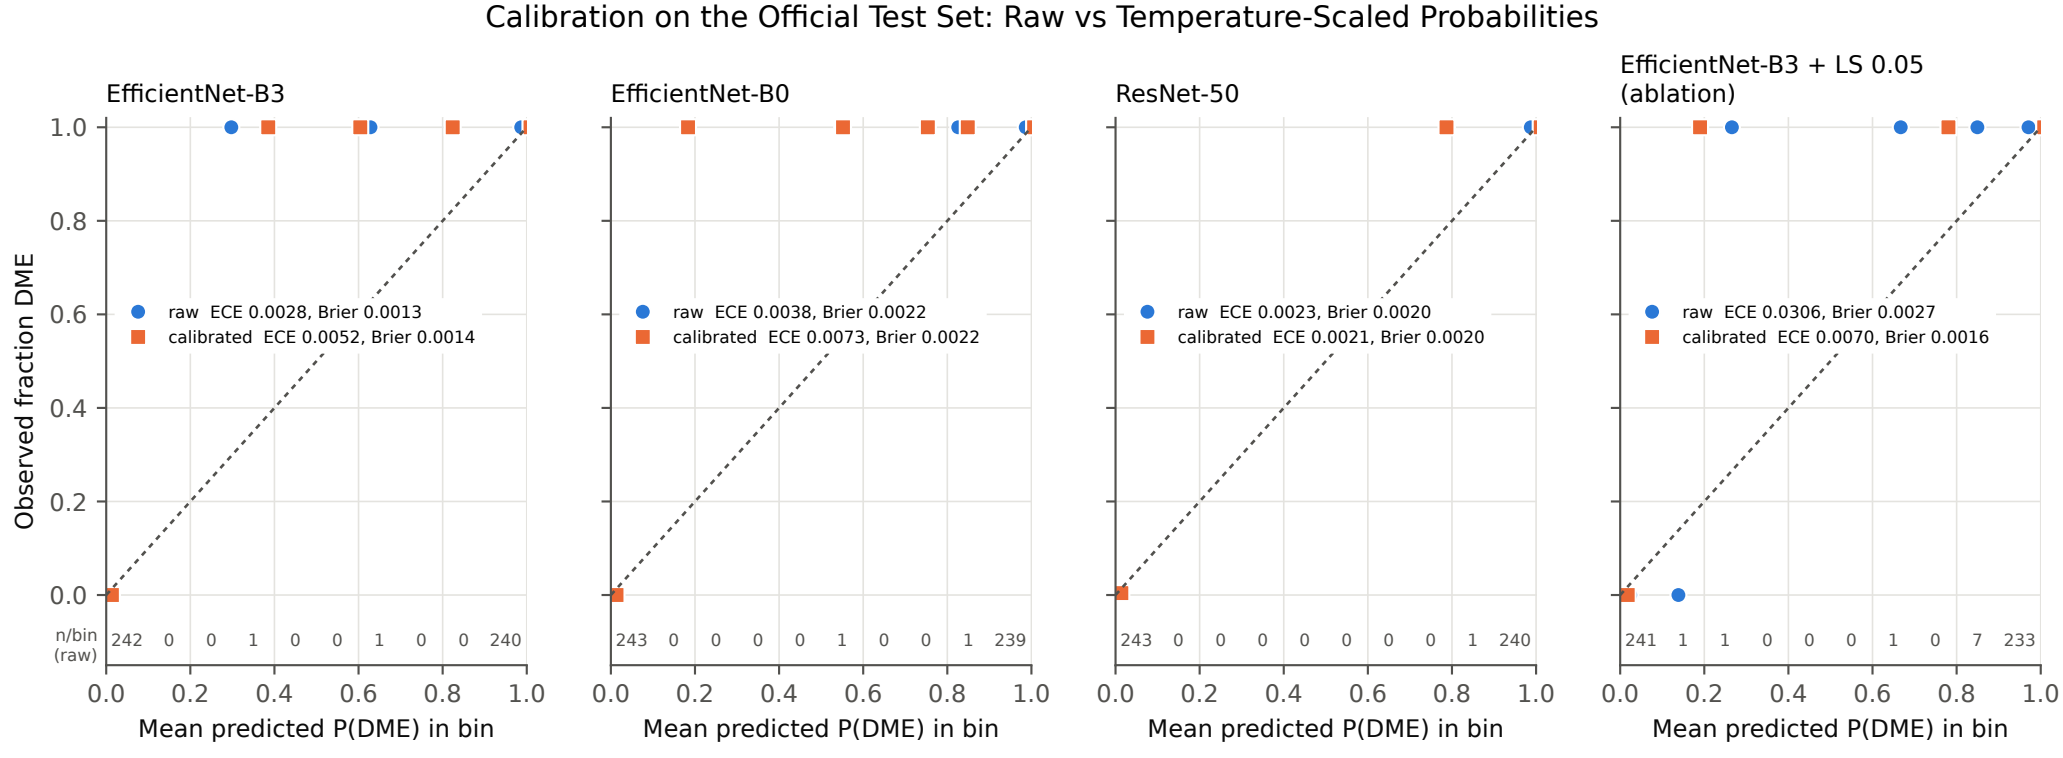

10 equal-width bins (same as evaluate.py); empty bins omitted; numbers under each axis = images per bin (raw). Almost all predictions fall in the outermost bins, so intermediate-bin estimates rest on 0–2 images.
Temperatures were refit with evaluate.fit_temperature on the training-time validation split (n=5,560, rebuilt with scikit-learn 1.4.2; scripts 11–12); the test set was not used for fitting.